


ACCRES Environments

SDT | 22 May 2018

What is a virtual environment?

- Isolated space with its own files, paths, libraries, dependencies
- The **conda environment** is a directory holding a set of installed packages

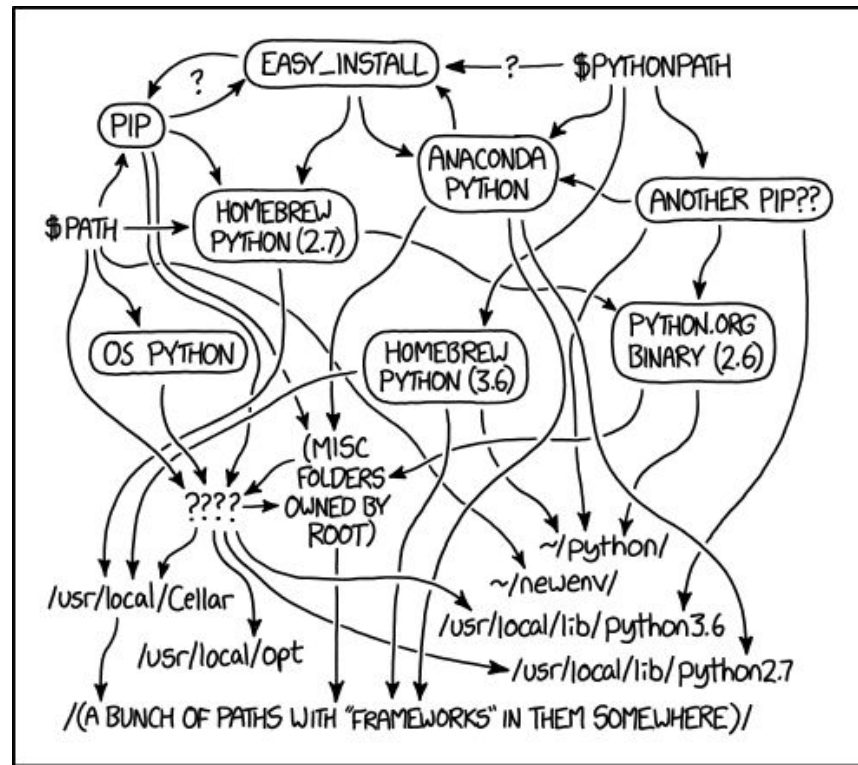
CONDA
ENVIRONMENTS
ARE MY
FAVORITE!



A true "python" environment | <https://medium.freecodecamp.org/why-you-need-python-environments-and-how-to-manage-them-with-conda-85f155f4353c>

Yes. You really need one.

- Manage complex dependencies
- Install new packages with ease
- Record all installed packages
- Run specific software/Python versions
- Simplify collaboration
- Avoid installation permission issues
- Maintain code reproducibility
- Isolate packages from each other & the global environment



MY PYTHON ENVIRONMENT HAS BECOME SO DEGRADED
THAT MY LAPTOP HAS BEEN DECLARED A SUPERFUND SITE.

Anaconda to the rescue

- Organized Python packages for data science and machine learning
- The **conda** package is a package, dependency and environment manager for any language



Lmod: the basics

Load Anaconda before creating/activating a conda environment:

```
module load Anaconda3
```

.....

KINDA LIKE
AN ACCRE
VIRTUAL ENV...

Save a specific module configuration with:

```
module save [ name ]
```

Load a saved module configuration with:

```
module restore [ name ]
```

ALSO:
MODULE Savelist
MODULE describe

How do I create an environment?

- Default create: **conda create --name gallifrey**
 - Environment will contain Python 2.7 if using Anaconda2
- Specify Python version: **conda create --name gallifrey python=3.5**

```
[colbrall@chgr1 ~]$ conda create --name gallifrey
Fetching package metadata .....
Solving package specifications:
Package plan for installation in environment /home/colbrall/.conda/envs/gallifrey:

Proceed ([y]/n)? y

#
# To activate this environment, use:
# > source activate gallifrey
#
# To deactivate this environment, use:
# > source deactivate gallifrey
#
```

Which environments do I have?

- `conda env list`

```
[colbrall@chgr1 ~]$ conda env list
# conda environments:
#
R3.4                /home/colbrall/.conda/envs/R3.4
gallifrey           /home/colbrall/.conda/envs/gallifrey
julia0.6.2          /home/colbrall/.conda/envs/julia0.6.2
root                * /opt/easybuild/software/Core/Anaconda2/4.4.0
```

- Alternatively: `ls .conda/envs/`

```
[colbrall@chgr1 ~]$ ls .conda/envs/
gallifrey  julia0.6.2  R3.4
[colbrall@chgr1 ~]$
```

Activating and Deactivating environments

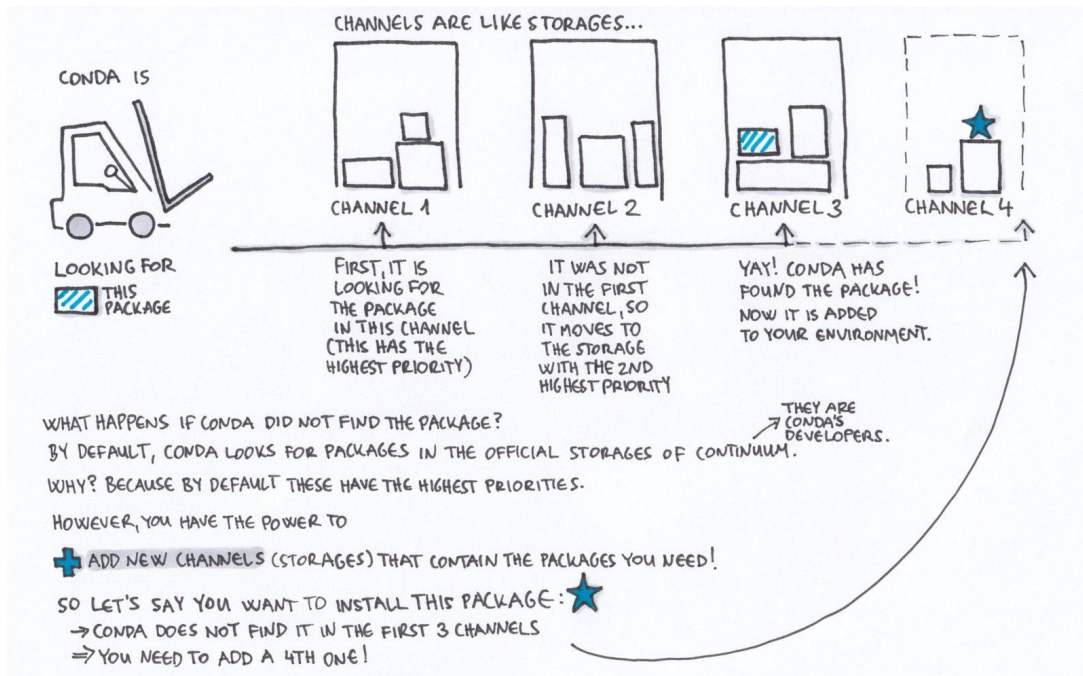
- **source activate gallifrey**

```
[colbrall@chgr1 ~]$ source activate gallifrey
(gallifrey) [colbrall@chgr1 ~]$ ls
00TEMPLATE.slurm  classes  data  intel  projects
(gallifrey) [colbrall@chgr1 ~]$ |
```

- **source deactivate gallifrey**

```
(gallifrey) [colbrall@chgr1 ~]$ source deactivate
[colbrall@chgr1 ~]$ ls
00TEMPLATE.slurm  classes  data  intel  projects
[colbrall@chgr1 ~]$ |
```


Installing/managing packages: channels



After conda 4.1.0, switch to channel priority > version.

We may change this behaviour with
`conda config --set channel_priority false`

Managing packages: channels

- `conda config --get channels`
- `Conda config --append channels
bioconda`
- `Conda config --prepend/--add
channels conda-forge`
- `Conda config --remove channels
conda-forge`

```
(gallifrey) chen111 ~ $  
conda config --get channels  
--add channels 'bioconda'    # lowest priority  
--add channels 'defaults'    # highest priority  
(gallifrey) chen111 ~ $  
conda config --append channels bioconda  
Warning: 'bioconda' already in 'channels' list, moving to the bottom  
(gallifrey) chen111 ~ $  
conda config --prepend channels conda-forge  
(gallifrey) chen111 ~ $  
conda config --get channels  
--add channels 'bioconda'    # lowest priority  
--add channels 'defaults'    # highest priority  
--add channels 'conda-forge' # highest priority  
(gallifrey) chen111 ~ $  
conda config --remove channels conda-forge  
(gallifrey) chen111 ~ $  
conda config --get channels  
--add channels 'bioconda'    # lowest priority  
--add channels 'defaults'    # highest priority
```

Managing packages: search

```
conda search [-c r] -f r-irkernel
```

```
(gallifrey) chenli1 ~ $  
conda search -c r -f r-irkernel  
Fetching package metadata .....  
r-irkernel  
0.2 r3.1.3_0 defaults  
0.2 r3.1.3_0a defaults  
0.2 r3.1.3_1 defaults  
0.2 r3.1.3_1a defaults  
0.3 r3.1.3_0 defaults  
0.3 r3.1.3_0a defaults  
0.4 r3.1.3_0 defaults  
0.4 r3.2.1_0 defaults  
0.4 r3.1.3_0a defaults  
0.4 r3.2.1_0a defaults  
0.4 r3.2.1_1 defaults  
0.4 r3.2.2_1 defaults  
0.4 r3.2.1_1a defaults  
0.4 r3.2.2_1a defaults  
0.5 r3.2.2_1 defaults  
0.5 r3.2.2_1a defaults  
0.5 r3.2.2_2 defaults  
0.6.1 r3.3.1_0 defaults  
0.7 r3.3.1_0 defaults  
0.7 r3.3.1_2 defaults  
0.7.1 r3.3.2_0 defaults  
0.7.1 r3.4.1_0 defaults  
* 0.8.9 r342hfe3cb8f_0 defaults  
0.8.11 mro343_0 defaults  
0.8.11 r343_0 defaults  
0.2 r3.1.3_0 r  
0.2 r3.1.3_0a r  
0.2 r3.1.3_1 r  
0.2 r3.1.3_1a r  
0.3 r3.1.3_0 r  
0.3 r3.1.3_0a r  
0.4 r3.1.3_0 r  
0.4 r3.2.1_0 r  
0.4 r3.1.3_0a r  
0.4 r3.2.1_0a r  
0.4 r3.2.1_1 r  
0.4 r3.2.2_1 r  
0.4 r3.2.1_1a r  
0.4 r3.2.2_1a r  
0.5 r3.2.2_1 r  
0.5 r3.2.2_1a r  
0.5 r3.2.2_2 r  
0.6.1 r3.3.1_0 r  
0.7 r3.3.1_0 r  
0.7 r3.3.1_2 r  
0.7.1 r3.3.2_0 r  
0.7.1 r3.4.1_0 r  
0.8.9 r342hfe3cb8f_0 r  
0.8.11 mro343_0 r  
0.8.11 r343_0 r
```

Managing packages: install and remove

conda install [-n gallifrey] [-c r] r-irkernel

```
(gallifrey) chen111 ~ $  
conda install -n gallifrey -c r -f r-irkernel  
Fetching package metadata .....  
Solving package specifications: .  
  
Package plan for installation in environment /scratch/chen111/conda_envs/gallifrey:  
  
The following NEW packages will be INSTALLED:  
  
  r-irkernel: 0.7.1-r3.4.1_0 r  
  
Proceed ([y]/n)? y
```

conda remove r-irkernel

```
(gallifrey) chen111 ~ $  
conda list | grep r-irkernel  
r-irkernel          0.7.1                r3.4.1_0    r  
(gallifrey) chen111 ~ $  
conda remove r-irkernel  
Fetching package metadata .....  
Solving package specifications: .  
  
Package plan for package removal in environment /scratch/chen111/conda_envs/gallifrey:  
  
The following packages will be REMOVED:  
  
  r-irkernel: 0.7.1-r3.4.1_0 r  
  
Proceed ([y]/n)? y  
  
(gallifrey) chen111 ~ $  
conda list | grep r-irkernel  
(gallifrey) chen111 ~ $
```

Share an environment

`conda env export > gallifrey_env.yml`

```
(gallifrey) chen111 ~ $  
conda env export > gallifrey_env.yml  
(gallifrey) chen111 ~ $  
more gallifrey_env.yml  
name: gallifrey  
channels:  
- defaults  
- bioconda  
dependencies:  
- _r-mutex=1.0.0=anacondar_1  
- backports=1.0=py27_0  
- backports_abc=0.5=py27_0  
- bleach=1.5.0=py27_0  
- bzip2=1.0.6=3
```

`conda env create -f environment.yml`

Delete an environment

```
conda env remove --name gallifrey
```

```
[colbrall@chgr1 ~]$ conda env remove --name gallifrey  
Remove all packages in environment /home/colbrall/.conda/envs/gallifrey:  
Proceed ([y]/n)? y  
[colbrall@chgr1 ~]$ conda env list  
# conda environments:  
#  
R3.4 /home/colbrall/.conda/envs/R3.4  
julia0.6.2 /home/colbrall/.conda/envs/julia0.6.2  
root * /opt/easybuild/software/Core/Anaconda2/4.4.0  
[colbrall@chgr1 ~]$ |
```

And don't forget to clean

Conda doesn't remove packages when you remove an environment, just the structure to access them.

conda clean --all removes any packages, caches, etc. that aren't used by any of your environments

Useful for staying under ACCRE file limits!

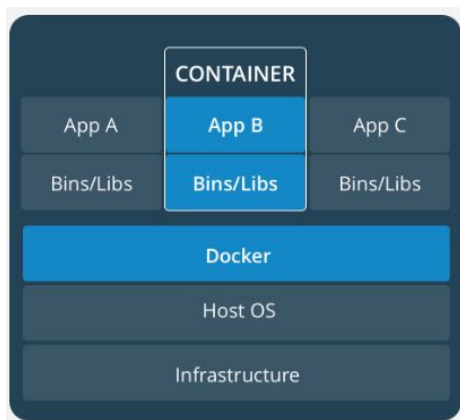
Conda is so powerful, why others?

Conda is good at managing packages for programming languages, but is not ideal to manage other dependencies.

Imagine we have to deploy a work pipeline with specific version of software, for instance, BEDTools/2.26.0. How can we ensure a reproducible, portable and immutable environment?

One solution is using a container.

What is a container?



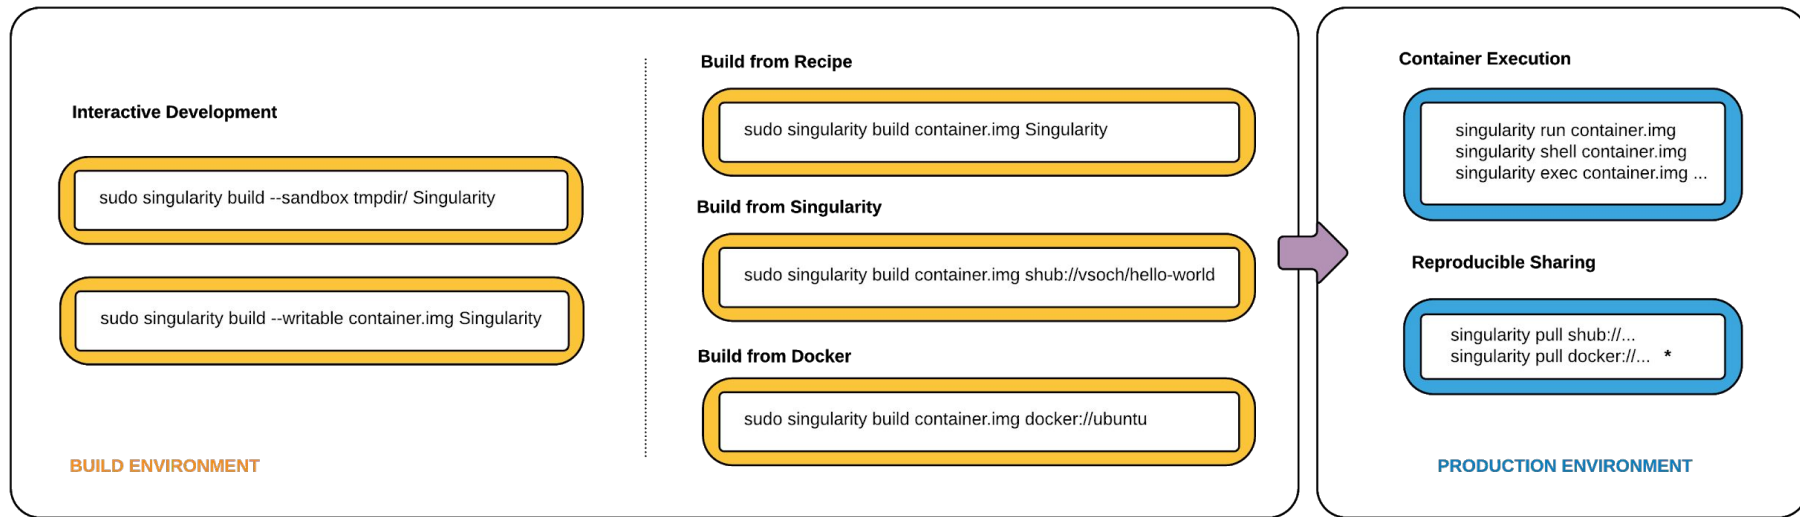
A container image is a lightweight, stand-alone, executable package of a piece of software that includes everything needed to run it: code, runtime, system tools, system libraries, settings.

A HPC container Singularity

Compare to docker, Singularity is designed for HPC computing. Because

- Security reasons. Docker gives root privileges. It's hard to limit access.
- Scheduling. It's hard to schedule resources to docker containers. On the other hand, slurm is built in to singularity.
- Singularity has better integration to *MPI*, GPUs...

The singularity workflow



* Docker construction from layers not guaranteed to replicate between pulls

The fastest way to use a singularity container

```
singularity shell shub://vsoch/hello-world
```

```
chen111 /dors/capra_lab/users/chen1/projects $  
singularity shell shub://vsoch/hello-world  
Progress |=====| 100.0%  
WARNING: Non existent bind point (directory) in container: '/scratch'  
WARNING: Non existent bind point (directory) in container: '/data'  
WARNING: Non existent bind point (directory) in container: '/dors'  
WARNING: Non existent bind point (directory) in container: '/gpfs22'  
WARNING: Non existent bind point (directory) in container: '/gpfs23'  
WARNING: Non existent 'bind path' source: '/legodata'  
Singularity: Invoking an interactive shell within container...  
  
Singularity vsoch-hello-world-master-latest.simg:~> whoami  
chen111  
Singularity vsoch-hello-world-master-latest.simg:~> █
```

A more advanced example: a singularity recipe

```
Bootstrap: docker
From: ubuntu

%help
Help me. I'm in the container.

%setup
touch ${SINGULARITY_ROOTFS}/tacos.txt
echo "I love avocados!" > avocados.txt

%files
avocados.txt

%labels
Maintainer Ling
Version v1.0

%environment
CAT=genius
export CAT

%post
mkdir /dors /scratch /data /gpfs22 /gpfs23

apt-get update
apt-get -y --force-yes install vim \
python-numpy \
python-pip && \
apt-get clean

%runscript
echo "Hello world from the container!"
echo "Argument received: $"
exec echo "$@"
```

Header: Define where to get the OS

Setup: Execute outside of container, after the base OS is installed.

Files: cp files into containers

Post: Execute inside of container at build time.
Meat of setup will live here. Install the packages you want here.

Runscript: Execute when the container is run, not at build time. Can take arguments.

Test it out!

```
vagrant@vagrant-ubuntu-trusty-64:~$ sudo singularity build test.img test.def
vagrant@vagrant-ubuntu-trusty-64:~$
vagrant@vagrant-ubuntu-trusty-64:~$ singularity help test.img

Help me. I'm in the container.

vagrant@vagrant-ubuntu-trusty-64:~$
vagrant@vagrant-ubuntu-trusty-64:~$
vagrant@vagrant-ubuntu-trusty-64:~$ singularity exec test.img cat avocados.txt
WARNING: Non existent bind point (file) in container: '/etc/localtime'
I love avocados!
vagrant@vagrant-ubuntu-trusty-64:~$
vagrant@vagrant-ubuntu-trusty-64:~$
vagrant@vagrant-ubuntu-trusty-64:~$ singularity exec test.img env | grep CAT
WARNING: Non existent bind point (file) in container: '/etc/localtime'
CAT=genius
vagrant@vagrant-ubuntu-trusty-64:~$
vagrant@vagrant-ubuntu-trusty-64:~$
vagrant@vagrant-ubuntu-trusty-64:~$ singularity run test.img you shake it all about
WARNING: Non existent bind point (file) in container: '/etc/localtime'
Hello world from the container!
Argument received: you shake it all about
you shake it all about
```

Useful Links

Anaconda site: <https://www.anaconda.com/download/>

Conda cheatsheet: https://conda.io/docs/_downloads/conda-cheatsheet.pdf

Conda-forge package list: <https://conda-forge.org/feedstocks/> (includes R, Julia packages)

Official Guide: <https://conda.io/docs/user-guide/getting-started.html>